

# Conversational Interface Design in the Age of LLMs

**Conversational interface design** sits at the intersection of human-computer interaction and the latest AI advances. It's about crafting the experience of talking *with* a machine (like a chatbot or voice assistant) in a natural, useful way. In this overview, we'll journey from the early days of chatbots to modern Large Language Model (LLM) systems like ChatGPT, and explore how interaction design and user experience (UX) principles apply to these AI-driven conversations. We'll also look at key design techniques (prompt design, system/user message framing, persona design, error handling, context management, feedback loops) and some creative examples of LLM-powered conversational interfaces – from AI therapists to storytelling game masters. Throughout, the tone will stay casual but professional, making these concepts accessible to an undergraduate reader.

## A Brief History of Conversational Interfaces

*A conversation with the first chatbot, ELIZA, created in 1966. ELIZA used simple pattern-matching techniques to imitate a psychotherapist's responses .*

Talking to machines is not a new idea – in fact, people have been experimenting with chatty computers for decades. The **first chatbot** was *ELIZA*, built in 1966 by MIT's Joseph Weizenbaum. *ELIZA* used pattern matching and scripted responses to mimic a Rogerian therapist, responding to user inputs with probing questions or reflections . This text-based program was very primitive by today's standards (as you can see in the screenshot above), but it laid the foundation for later conversation programs.

After *ELIZA*, there were many other early chatbots and conversational agents. In the 1990s, for example, *ALICE* (1995) used a set of heuristic pattern rules to chat, and in the early 2000s the MSN/AIM bot *SmarterChild* could banter and fetch simple info – a precursor to the smarter assistants to come . The 2010s then saw the rise of mainstream **voice assistants** like Apple's Siri (launched in 2010) and Amazon's Alexa (2014). Siri and Alexa could understand spoken requests and respond via speech, using natural language processing to turn commands like "set a timer for 10 minutes" or "play some music" into actions .

These voice-driven interfaces made conversational tech familiar to millions of people.

Despite these advances, for a long time chatbots remained fairly limited in capability and scope. They were mostly used for **simple, scripted tasks** – answering FAQs, setting reminders, or basic customer support – and they often struggled if users went off the “happy path” of expected inputs. Until recently, many users found chatbots more frustrating than useful; one report in 2022 noted that *72% of people found chatbots to be a waste of time* . That perception flipped almost overnight with the emergence of modern LLM-based bots. In late 2022, **OpenAI’s ChatGPT** burst onto the scene and gained over 100 million users within a few months . Unlike the old rule-based bots, ChatGPT could engage in free-form conversations on almost any topic, producing responses that (mostly) felt coherent and contextual. This breakthrough was powered by large-scale neural network models (transformers) trained on massive text datasets – but from a user’s perspective, it felt like chatting with a really knowledgeable, talkative person. The sudden popularity of ChatGPT and similar models in 2023–2024 marks a new chapter in conversational interface design, where *conversation itself is becoming a mainstream interface* for all sorts of applications .

## From HCI to Conversation Design: Interaction Design Principles

Even with advanced AI under the hood, a good conversational system doesn’t happen by accident – it requires thoughtful **conversation design**. Conversation design is often defined as *“the practice of making AI assistants more helpful and natural when they talk to humans,”* combining technology, psychology, and language insights to create human-centric experiences . In other words, we apply the same thinking from general **interaction design and UX** – understanding users, designing clear interactions, ensuring usability – but in the context of a back-and-forth dialogue instead of a graphical interface of buttons and links.

Some fundamental principles from interaction design and UX **directly apply to conversational UIs** (and are emphasized by conversation-design experts):

- **User-Centric Design:** Start with the user’s needs and goals. A conversational agent should solve real problems or provide real value for the user in an intuitive way . This means anticipating what the user is trying

to do and making it as easy as having a natural conversation. For example, if users want to book a flight via chat, the flow of questions and answers should feel logical and not require the user to struggle or repeat themselves unnecessarily.

- **Make it Helpful and Cooperative:** In human conversation, we expect a cooperative partner. Similarly, an AI assistant should actively try to be helpful, relevant, and clear. This principle underlies everything from how the bot interprets user input (e.g. inferring intent even if wording isn't exact) to how it formulates responses (being concise and on-topic). Conversation design draws on linguistics maxims and psychology so the AI follows norms like not lying (quality), not rambling (quantity), staying relevant, and being clear in its wording – all contributing to an overall useful interaction.
- **Maintain Context:** One big part of making conversation feel natural is *context awareness*. The system should remember what has been said earlier and use it to inform later responses . If you ask a follow-up question like "How about tomorrow?" after asking about today's weather, the AI should understand you're still talking about the weather in the same location, not suddenly lose track . Preserving context avoids forcing the user to repeat or rephrase things that should be "understood" from prior turns.
- **Consistent Persona and Tone:** Just as visual interfaces follow a style guide, conversational interfaces should have a consistent voice and personality that suits the use case. Is the assistant formal and professional, or friendly and playful? This **persona design** should match user expectations and the brand's character . For instance, a banking assistant might be polite and serious to convey trust, whereas a teen-focused chatbot might use more casual language and emojis. The key is consistency – the AI shouldn't sound like a goofy friend in one reply and a stodgy lawyer in the next (unless deliberately designed to switch personas). A well-defined persona builds familiarity and trust over time.
- **Error Handling & Recovery:** In any interaction (GUI or conversation), things can go wrong or get misunderstood. Good conversational UX means planning for errors, misunderstandings, or "no match" situations and handling them gracefully . If the AI doesn't understand the user's request or if something goes off track, it shouldn't dead-end or give a gibberish answer. Instead, it might apologize and ask a clarifying question ("I'm sorry, I didn't catch that. Could you rephrase?"), or provide a helpful fallback ("I'm

not sure about that – I can help with flight bookings or weather info, for example.”). The design goal is to *recover* the conversation or at least fail gracefully without frustrating the user. Limit how many errors in a row the user endures – after a couple of misfires, it might offer to hand off to a human agent or provide another way to get the answer .

- **Guidance and Flow Control:** In traditional interaction design, we guide users with visual cues; in conversation, guidance happens through prompts and the structure of dialogue. A well-designed chatbot uses prompts or questions to steer the interaction helpfully. It might provide options (“Are you looking to track an order, return an item, or something else?”) to make it easier for the user to answer. This ties into **dialogue management** – controlling the flow so it’s neither too open (which can confuse the AI or user) nor too rigid. With LLMs that can handle free-form input, designers often still impose some structure or hints to keep the chat productive.
- **Channel Appropriateness:** If the conversation is happening via voice (smart speaker) versus text (chat window), the design should respect the medium. Voice interfaces may need shorter responses (since listening is slower than reading) and more auditory cues. Text chat can use visual elements (like quick-reply buttons, images, or formatting) if the platform allows. Great conversational UX considers these context differences – for example, a voice assistant should confirm important info back to you (“Okay, booking your flight to Denver on June 5th...”) because mishearing could be costly, whereas in a text UI the confirmation might be in a neatly formatted summary. Tailor the conversation style to the channel’s strengths and constraints .
- **Continuous Improvement:** The design process doesn’t end at launch. Real conversations with users will reveal new needs, quirks, and failure cases. **Conversation designers iterate** on the dialog flows and prompts by reviewing chat transcripts, user feedback, and analytics . This is akin to usability testing for GUIs. If users keep getting confused at a certain question, redesign it; if the bot is frequently hallucinating information, adjust the prompts or constraints. Many teams incorporate feedback loops where the AI might even ask the user if its answer was helpful, providing an explicit rating mechanism. Over time, these improvements make the experience more robust and delightful.

In summary, the fundamentals of **good UX** – being user-focused, keeping things intuitive, providing feedback, handling errors – all carry over to

conversational interfaces. The twist is that instead of clicking and tapping, the interaction is through language. So, principles from linguistics and human conversation etiquette matter as much as GUI design rules. Next, we'll delve into specific techniques that designers and developers use to achieve these goals with LLM-based systems.

## Key Design Techniques for LLM-Powered Conversations

Modern LLMs like GPT-4 are powerful but also unpredictable. Designing a usable conversational experience with such models involves various *techniques and tricks* to align the AI's behavior with user needs. Here are some of the most important techniques, and how they relate to interaction design:

### 1. Prompt Design and Engineering

Because an LLM doesn't have a fixed, hard-coded behavior, **how you prompt it** largely determines what it will do. *Prompt design* (aka prompt engineering) is the art/science of phrasing the input or instructions to the model to get the desired output. As one expert puts it, *you can't program an LLM in a traditional sense; "you need a vision of what you want to achieve and mimic the initiation of that vision" – this process of mimicking is prompt design*. In practice, that means when we design a conversational turn, we carefully craft the text we send to the model (which may include the conversation history, system instructions, and the user's last question) to guide its response.

Some key considerations in prompt design for conversations with LLMs:

- **Providing Context:** LLMs are essentially stateless in that they only know what's in the prompt (plus their training data). To make them perform well, we often need to include relevant context in the prompt. For example, if a user asks a follow-up question, the prompt should usually include the previous Q&A turns so the model has that context. If we want the AI to answer based on a specific document or data, we need to put that text into the prompt (a technique known as *context stuffing* or retrieval-augmented generation, where the prompt includes snippets of relevant info). A good prompt ensures the model isn't answering in a vacuum.
- **Clarity of Request:** A well-designed prompt clearly indicates what the user (or system) is asking for. Ambiguous prompts lead to meandering or incorrect answers. Research into how users write prompts for AI has found

that the best prompts often have a structure: they include a direct request or question, some framing or context, maybe an example or reference, and sometimes a desired format for the answer. For instance, instead of just saying “Tell me about Mars”, a user might have more success with “I’m doing a school project. **Give me a brief summary** of key facts about Mars, **in bullet points**.” This prompt explicitly states the task (summary of key facts) and format (bullet points), which helps the LLM focus and produce a useful response.

- **Prompt Structure & Examples:** Designers sometimes use *few-shot prompting*, which means including examples of Q&A or dialogue in the prompt to shape the model’s style and knowledge. For example, to teach the model how to respond as a friendly tutor, a prompt might literally show a mini conversation between a student and tutor, and then ask the model to continue in that vein. Structured prompts can also mitigate issues like hallucinations. One approach is giving the model a template to fill in – e.g., “Context: [some text]\nQuestion: [user question]\nAnswer:”. This can anchor the model to use the provided context and say “I don’t know” if the answer isn’t in the context. Essentially, by structuring the prompt, we guide the model’s reasoning process and format its output.
- **Iterative Refinement:** Often, prompt design is an iterative process. You try a wording, see how the model responds, then adjust. Even slight rephrasing or adding a detail can change the output. This is why it’s called “engineering” – you experiment and refine the input to engineer the desired outcome. For instance, if an AI keeps giving too long an answer, you might add “Answer in 3 sentences.” If it speaks in first person when you want third person, you explicitly instruct the style in the prompt. Over time, designers accumulate prompt patterns that work well (there are even public “prompt libraries” sharing tricks to get models to behave in certain ways).

In short, prompt design is crucial because it’s how we **communicate our design intentions to the model**. A well-designed prompt can make an LLM seem incredibly smart and helpful; a poorly designed one can lead to confusion or nonsense. As a conversational UX designer, you often have to **think about phrasing** and what information to include or omit in each turn’s prompt.

## 2. System vs User Messages (Framing the Conversation)

Many LLM-driven chat systems (like OpenAI’s API, ChatGPT, etc.) distinguish between **user messages** and **system messages** (and sometimes **assistant**

**messages**). This is essentially a mechanism to give the AI some background instructions *separately* from the user's query. The **system message** is an invisible guiding prompt that sets the stage for the conversation – it's where designers can inject context, rules, or persona from the get-go . The **user message** is what the human actually says/asks at each turn, and the **assistant message** is the AI's reply.

Using the system message effectively is a key design technique:

- **Defining the Role and Persona:** The system message often contains a short description of *who or what the assistant is*. For example: *"You are a helpful travel agent bot who speaks in a friendly, concise manner."* By giving the model a role, we anchor its style and behavior. OpenAI notes that a system message can include the assistant's personality traits or a brief backstory . This is essentially **persona design via prompt**. If we want a playful tone, we might say *"You are a witty, creative storyteller AI..."*; for a customer service bot, *"You are a polite and professional virtual assistant for a bank..."*.
- **Setting Instructions and Boundaries:** The system message is also where we can put **conversation ground rules**. For instance, we could instruct: *"Only answer questions about cooking. If the user asks something outside of cooking, politely refuse."* or *"If you don't know the answer, say 'I'm not sure' instead of making something up."* These act as guardrails to prevent unwanted behavior or off-topic answers . In UX terms, it's like configuring the AI's constraints so it doesn't violate user expectations or company policy. The user doesn't see these instructions, but the AI does and will (hopefully) follow them.
- **Providing Contextual Knowledge:** We can also use the system message to feed in context that should persist throughout the chat. If a chatbot has access to, say, an FAQ or knowledge base, a designer might stuff some of that info into the system prompt (within token limits) or use it to remind the AI of facts. For example: *"Knowledge: The company was founded in 2010 and now has over 500 stores..."* followed by *"You are an assistant who answers questions about this company."* This way the AI always has that background while conversing .
- **Overall "Atmosphere" of the Chat:** The system message helps shape the general *mood*. If we want a casual, welcoming chat, we might include language like *"The assistant is friendly and upbeat."* If it's a more serious

domain, we reinforce professionalism. This is subtle but it does affect how the model responds over long conversations.

Using system messages is a bit like the **"setup"** before the user even says anything – it frames the interaction. Meanwhile, the **user messages** are just whatever the user inputs. As conversation designers, we can't change what users will say, but we can sometimes influence it (through our prior prompts or by offering suggested replies in a UI). However, we *can* structure how we feed user messages + system instructions to the model. By carefully designing that combination, we create a more reliable and on-brand conversational experience. In effect, the system message acts as an "invisible moderator" of the conversation: it's always there pushing the AI toward certain behavior.

### 3. Persona and Tone Design

As mentioned earlier, defining a **persona** for the AI is important for consistency and user engagement. Persona design in conversational interfaces involves deciding *who the bot is (from the user's perspective)* and how it speaks. This can include a name, a speaking style, a level of formality, even a bit of implied personality. The goal isn't to fool the user that it's human, but to give the interaction a coherent character – something we can relate to, or at least **find predictable and appropriate**.

Techniques for persona design with LLMs:

- **System Message Persona Script:** The primary method is to encode the persona in the system prompt (as discussed). For example, if we're building "TutorBot", our system message might say: *"You are TutorBot, a friendly online tutor who helps high school students with math homework. You are patient and use simple language. You crack an occasional mild joke to put students at ease."* This one-time instruction can profoundly shape the tone of every answer the AI gives, making sure it doesn't, say, lapse into super dry encyclopedia mode. A concrete example: one team designed a bookstore assistant persona with a system prompt *"You are a helpful assistant for a bookstore called BookBuddy. Give short, courteous answers, no more than 1 sentence. Always be accurate. If you don't know the answer, say so."* – this ensured the AI's replies stayed concise, polite, and on-topic .
- **Style Guides and Example Dialogues:** Designers might create a **style guide** for the chatbot's language, similar to branding guidelines. This could list preferred words or phrases (e.g. calls user "guest" vs "customer", or uses



contractions to sound casual), as well as things to avoid (no slang, or conversely, no overly technical jargon). Feeding a few example dialogue snippets that embody this style can help an LLM pick up the voice. For instance, showing how the bot would introduce itself to a user in that persona.

- **Consistent Behavior Patterns:** Persona is not just about tone, but also behavior in conversation. A therapy bot persona might consistently respond with empathy and questions (as a human therapist would), whereas a tech support bot persona might be more transactional and step-by-step. These patterns can be part of the design. An AI therapist could be instructed in the system prompt to “always validate the user’s feelings and ask an open-ended question about their situation,” reflecting a known counseling technique. A troubleshooting bot might be instructed “if the user is frustrated, apologize for the inconvenience first, then suggest a next step.” By building these behaviors into initial prompts or into the conversation flow logic, the *user experiences a reliable character* on the other end.
- **Adapt Persona to Audience:** Just like any product should know its audience, a bot’s persona should resonate with its users. If the target users are kids, the persona might be playful and metaphor-rich. If the users are doctors using a medical AI assistant, the persona should be knowledgeable and use the proper terminology (while still being courteous). A key part of persona design is understanding the users’ expectations – we wouldn’t want Clippy-like jokes in a serious medical advice context, for instance. Thus, persona design often starts with questions like “*Who will use this? What tone would they appreciate or trust?*” .
- **Humanizing (within reason):** There’s evidence that users respond well when a bot has a dash of personality or humor *as long as it remains helpful*. For example, a weather bot that says, “Don’t forget your umbrella 🌧️” might delight users more than one that says “Rain likely. 90%.” – even though the latter is perfectly factual, the former feels like a friendly nudge. The persona we give the AI can include these little touches. However, designers must be cautious: overdoing the persona (especially pretending it has emotions or using too much slang) can backfire and seem creepy or annoying. It’s a balancing act to keep the experience natural and user-centric.

In summary, **persona design** ensures that an LLM-powered assistant isn’t just a raw information spitter, but a *designed character* that makes the interaction

more engaging and suitable for the context. By controlling persona, we indirectly help with consistency (the bot won't suddenly shift tone) and set user expectations (they get a sense of "who" they're chatting with). This ultimately contributes to a better UX because people know what kind of responses to expect from a given persona.

## 4. Error Recovery and Graceful Failure

No matter how well we design prompts and personas, conversations with AI will sometimes go off the rails. The AI might misunderstand the user's question, the user might type something incoherent, or the model might just *make a mistake* (e.g. giving wrong info or "hallucinating" a non-existent fact). A critical part of conversational interface design is planning for these moments and handling them in a user-friendly way.

Approaches to error handling in conversational design:

- **Fallback Responses:** These are the bot's equivalent of "error messages," but phrased in a conversational manner. For example, if the AI's confidence is low or it's unsure how to handle something, it could use a fallback like *"I'm sorry, I'm not sure about that. Maybe try rephrasing your question?"* or *"I don't have that information, unfortunately."* It's important that these fallbacks **don't blame the user** (never say "You are not making sense" or such) but rather apologize or express the limitation as the bot's own. The tone should remain polite and helpful. Designers often script a few of these and instruct the AI (via system prompt or logic) to use them when needed.
- **Clarification and Confirmation:** Sometimes the best way to handle a possible misunderstanding is to ask the user for clarification. If a user says *"Book ticket for Friday"* and the AI isn't sure which event or destination, a good recovery is to ask a follow-up: *"Sure, I'd be happy to help – did you mean a movie ticket, or a flight, or something else for Friday?"* This turns a potential error into an *active query* that guides the user to provide the missing info. Similarly, if the AI isn't certain it heard correctly (in voice UIs) or parsed correctly, it might confirm: *"Just to confirm, you'd like a ticket for this upcoming Friday, is that right?"* This technique of confirming critical info helps catch misunderstandings early. It's analogous to a form validating input: better to double-check now than proceed on wrong info.
- **Limiting Repetition:** If an error happens repeatedly, users will get frustrated quickly. A common design rule (from voice UI design guidelines) is the "three-strikes rule" – if the assistant fails to understand after two attempts,

on the third it should have a strategy to escalate or give up gracefully . For instance, after two “Sorry, could you repeat that?” loops, the third time it might say *“I’m having trouble understanding. Let me connect you to a human agent.”* or *“Maybe try using our website for this request – here’s the link.”* This prevents infinite loops of confusion.

- **Preventing Known Pitfalls:** Through testing, designers may identify certain questions or phrasing that consistently confuse the AI or lead to bad answers. If those are predictable (like a math word problem that the model fails at, or a trigger phrase that causes it to ramble off-topic), designers can create **rules or additional prompt instructions** to handle them. For example, if an AI for a bank often gets asked about something it *can’t* answer (say, detailed investment advice), the system prompt might include: *“If user asks about investments, respond with a disclaimer and suggest scheduling a call with a financial advisor.”* This way, even if the model doesn’t truly “know,” it has a safe, pre-designed response. Essentially, we patch the holes we know about.
- **Using the Model’s Uncertainty (Advanced):** Some advanced techniques involve checking the AI’s “uncertainty” in its answer (if the platform provides a way to do that) or using multiple phrasing to see if it’s consistent. While not foolproof, designers might use heuristics – e.g., if the AI’s answer contains a phrase like “I’m not sure” or seems to contradict itself, treat it as low-confidence. Then trigger a different prompt or response. Another approach is having the AI explain its reasoning internally (chain-of-thought prompting) and verify it. These are more implementation-level details, but they tie into design: the goal is to catch errors before they reach the user, and either fix them or present a controlled failure message.

The overarching UX principle is **graceful degradation** – when things go wrong, the experience should *bend, not break*. The conversation might deviate from the happy path, but the system guides it back or at least communicates limitations clearly. By designing and incorporating error recovery strategies, we make the AI feel more reliable and trustworthy. Users will be forgiving of an error if it’s handled gracefully, but a bad response that’s confusing or incorrect with no acknowledgment can erode confidence fast. So this area is absolutely crucial for maintaining a good experience.

## 5. Context Management in Dialogue

Context is the lifeblood of conversation. Humans keep track of context effortlessly – pronouns refer to earlier nouns, implied topics carry over turns, etc. LLMs have some capacity to do this (thanks to their sequence-based input), but they have a finite memory (context window) and can lose track over a long chat. Designers need to employ tactics to **manage context** so that the interaction remains coherent and relevant even as the conversation grows.

Key aspects of context management:

- **Referencing Previous Turns:** The simplest form is just including the recent conversation history when calling the model. Most chat interfaces do this automatically: when the user sends a new message, the prompt to the LLM includes some or all of the prior messages so far. However, there's a limit (e.g. maybe the last 10-20 messages if the model's context window is limited). Designers might decide how much history to include – too little and the bot forgets context, too much and it might exceed limits or get confused. One strategy is **summary compression**: summarize older parts of the chat and include the summary instead of every line. This keeps important facts in context without the full detail. For example, if earlier the user provided their name and problem description, and now we're 30 minutes into troubleshooting, the system might carry a summary like "(User is John, having issue with phone not charging)" in the prompt.
- **Slot Filling & Memory:** In task-oriented dialogues (like booking or forms), context can be managed through explicit slots. E.g., the bot collects Date: \_\_, Time: \_\_, Location: \_\_ and once filled, it *knows* those values going forward. Even if not using a formal slot architecture, the conversation design should ensure once information is given, it's acknowledged and reused. If you already told an AI your name is Alice, it should ideally address you as Alice later or not ask again needlessly. LLMs don't have long-term memory across sessions by default, but conversation designers can implement **user profiles** or database lookups. For instance, if a returning user says "What's my account balance?", the system might retrieve their identity from a login and then let the AI know (via the prompt) "User is Mike and his account balance is X as of June 1." This mixes programmed context with AI, creating a more personalized continuous experience.
- **Context Switching and Coreference:** A robust conversational interface can handle when the user changes topic or refers to something ambiguously. If a user asks multiple questions in one turn or jumps to a new subject, the

design should accommodate that rather than break. Some LLMs can handle multiple queries in one prompt (maybe answering each in turn). If not, the bot might need to gently manage it: *"I'd be happy to help with that too, but let's finish your first question about your schedule, then we can address the new question about budget."* Coreference (pronouns like "it", "they") is another challenge – the AI should infer what "it" refers to based on context. For example: *"I met Mark and Tom yesterday. He was excited about the project."* A good AI should know "he" likely means Mark (depending on context clues). While we rely on the model's language ability for this, designers can help by how they present information. Sometimes rephrasing user input or injecting explicit names can prevent confusion (though doing so must be done carefully to not distort what user said). In multi-turn interactions, if there's ambiguity, the bot can ask: *"Just to clarify, when you said 'he', did you mean Mark?"* This ties back to error handling and clarification strategies.

- **Context Limits and Resets:** With very long conversations, an LLM might eventually hit its limit or start to get things wrong due to earlier information being out-of-sight (forgotten) or contradictory info introduced. In some cases, starting a "new session" or summarizing and clearing is needed. A chatbot might say, after a long tangent, *"Anyway, how else can I help you today?"* which subtly resets context to a new topic. Another tactic: if the conversation goal is achieved (ticket booked, issue resolved), the bot can close out and suggest starting fresh for new inquiries. This prevents old context from bleeding into a separate task. It's a bit like clearing a form after submission – clear the conversation context when appropriate.

In essence, context management is about **keeping the AI on the same page as the user**. It's part technical (how much we feed into the model) and part design (how we prompt, clarify, and possibly constrain the dialogue). When done well, the user feels like the AI "remembers" everything they've said – which makes the experience feel much more natural and intelligent. When done poorly, the AI might ask for info twice, or change the subject unpredictably, or contradict itself, breaking the illusion of a coherent conversational partner.

## 6. Feedback Loops and Continuous Improvement

Designing conversational experiences with LLMs is not a one-and-done effort – it requires **ongoing tuning and learning from interactions**. There are two kinds of feedback loops to consider: **real-time feedback within the conversation**

(where the user's reactions guide the immediate next steps) and **post-conversation feedback for designers/developers** (to improve the system over time).

Real-time interactive feedback techniques include:

- **Implicit Feedback:** In conversation, if the AI's answer is unsatisfactory, the user will often indicate that in their next response (e.g. *"No, that's not what I meant."* or *"That's incorrect."*). A well-designed system should recognize these cues and treat them as feedback to adjust. For instance, if the user says *"No, that's wrong,"* the bot can reply *"Oops, sorry about that. Let me try again – could you clarify what part I got wrong?"*. The conversation design should make the AI *flexible to course-correct*. LLMs are capable of taking new instructions, so the prompt could incorporate the user's negative feedback: e.g., *"The user said that was wrong. Assistant, revise your answer considering that..."*. This way, the interaction itself becomes a feedback loop where the user teaches the AI what they need. It's very much like how a human conversation works (*"Actually, let me rephrase..."*).
- **Explicit Feedback Requests:** Some chatbot interfaces include options for the user to thumbs-up or thumbs-down a response, or to answer *"Was this helpful? (Yes/No)"*. If the user says it wasn't helpful, the system might either adapt immediately (offer an alternative answer) or at least apologize and try a different approach. Designers have to be careful not to annoy the user with too many confirmation questions, but well-timed feedback prompts can improve the experience. For example, after a long explanation, the bot might ask *"Did that answer your question?"* – if the user says no, the bot can attempt a clarification or escalate. This immediate feedback loop ensures the user's need is ultimately met.
- **Adaptive Response based on Sentiment:** More advanced systems might analyze user sentiment or tone during the chat. If a user is getting frustrated (maybe their messages start looking angry or they've asked the same thing multiple times), the AI could adapt by switching to a more sympathetic tone, or simplifying its language, or quicker suggestion to get human help. This is similar to how a customer service agent might handle an irate caller by calmly acknowledging the frustration and trying a different tactic. By monitoring the "feedback" inherent in *how* the user is speaking, an AI can modify its strategy – a form of real-time UX optimization.

Long-term feedback loops involve learning from all conversations to improve the system:

- **Analyzing Chat Logs:** Conversation designers or analysts will review transcripts of chats (with privacy in mind, of course) to spot common failure points or opportunities. For example, they might find that many users ask a certain question that the bot doesn't handle well – that's a cue to add information or an explicit rule for it. Or they might notice users frequently using a phrase the AI misinterprets; the designers can then add that phrasing (and its intended interpretation) to the training data or adjust the prompt to handle it. Essentially, real user chats are goldmines for understanding what users *actually* want and how the AI is falling short.
- **User Ratings and Surveys:** If the platform collects user ratings (like those thumbs up/down), that data can feed into model improvement. A very direct way is using those ratings to fine-tune the model (this is part of the Reinforcement Learning from Human Feedback – RLHF – approach that was used to train ChatGPT). Even without full RLHF pipeline, designers can look at, say, conversations with many thumbs-down and diagnose what went wrong. Some systems also ask a brief survey at the end of a chat ("How was your experience?"). All this helps identify satisfaction and pain points.
- **Iterative Prompt and Policy Updates:** With insights from logs and feedback, the conversation design is refined. This might mean updating the system message instructions, adding new example prompts, tweaking fallback responses, expanding the knowledge base, or even adjusting the model if possible. It's an iterative cycle: design → deploy → collect feedback → redesign. A mindset of **continuous improvement** is crucial because language interactions are so dynamic – users will always surprise you with new ways to ask things, or the world changes and the bot needs updating (e.g., new product names, new slang terms, etc.). Companies often schedule regular training updates for their models or roll out dialogue improvements on a rolling basis as they learn.
- **Learning from Mistakes at Scale:** One advantage of AI systems is that every mistake can, in theory, be analyzed and learned from globally. Unlike a human agent who might individually learn, here the *system as a whole* can be improved so the mistake ideally never happens again with any user. For example, if an AI assistant in a game gave a bizarre response in a particular scenario, developers can patch that. Or if a medical advice bot gave an unsafe suggestion, the team can add a rule to prevent such

suggestions in the future (and update the model). Over time, these corrections make the AI more robust. Essentially, the AI's design is *never static*. As one blog noted, conversation design is a continuous process of refining the experience using real data from chats .

By incorporating feedback loops, we ensure that the **"user experience design"** of the AI gets better over time, much like how apps issue usability updates. It also fosters user trust: if people see the AI improving or quickly correcting itself when guided, they feel more in control and heard. In an educational setting, for example, if a student says "I don't understand that explanation," and the tutor bot then tries a different way to explain (simpler language or a metaphor), the student feels like the system is responsive and caring. That's the UX win that feedback loops enable.

## Creative and Surprising Applications of LLM Conversations

Designing with LLMs opens up a playground of new and unexpected use cases for conversational interfaces. Beyond the typical customer service chatbot or virtual assistant, we're seeing LLM-driven agents in roles that would have seemed like science fiction just a few years ago. To spur your imagination, here are a few **creative examples of conversational LLM applications** and how design considerations come into play for each:

- **AI Therapy and Mental Health Support:** Therapy chatbots have existed for a while (recall ELIZA was essentially an early psychotherapy bot!), but LLMs have made them far more realistic and responsive. Chatbots like *Woebot* and *Wysa* guide users through Cognitive Behavioral Therapy exercises or just have empathetic conversations, helping people with anxiety or depression. Studies have found that such bots can **improve users' depression symptoms and build a therapeutic alliance** similar to that of human therapists . Designing an AI therapist involves a careful persona – it must be extremely empathetic, patient, and never judgmental. The prompts need to encourage the user to open up, while the AI must be constrained from giving harmful advice. Techniques like **tone matching** (mirroring the user's emotional tone in a supportive way) are used. For example, if a user expresses sadness, the bot might respond, *"I'm sorry you're feeling that way. That sounds really tough."* – showing empathy before perhaps suggesting a coping exercise. Error handling here is critical: if the AI doesn't understand what the user said (maybe the user is in crisis or uses



slang the AI missed), it should not produce a non-sequitur. Many systems have safeties: keywords that trigger an escalation (like providing the number for a crisis hotline or alerting a human moderator if someone mentions self-harm). Privacy and trust are also huge – users need reassurance their confessions aren't going to be tweeted out by a rogue bot. So, designers explicitly mention confidentiality and use secure channels. An LLM can feel surprisingly *human* in these settings, sometimes giving people an outlet to talk about things they might hesitate to tell a human. It's an exciting area, but one requiring ethical caution and robust design to truly help, not harm.

- **Storytelling and Creative Writing Companions:** LLMs can spin stories on the fly, which has given rise to interactive storytelling bots. For instance, **AI Dungeon** is a text-based adventure game where the player and an AI collaboratively tell a story – the AI generates imaginative scenarios and responds to the player's actions in a game world . Similarly, there are writing assistants that help authors brainstorm plots or write in the style of famous authors. Designing a storytelling AI means embracing the model's creativity while guiding it to maintain coherence. As a designer, you'd ensure the prompt always includes the *current story state* (characters, setting, etc.) so the AI doesn't contradict itself. You might give it instructions like *"keep the narrative in second person present tense, and don't abruptly change the character identities."* Users often want to be surprised by the AI's ideas, so you allow a certain randomness (higher "temperature" setting in AI terms) – this can lead to delightfully unexpected plot twists. However, the UX should allow the *user* to steer the story. Some interfaces let users undo or redirect if the AI's contribution wasn't to their liking (for example, "try a different action" if the story went in a weird direction). A well-known challenge is the AI's tendency for **"algorithmic absurdity"** – it might suddenly output something nonsensical or break the story's logic . Designers mitigate this with a mix of prompt engineering (keeping it grounded in the story context) and giving users control to correct the course. The result, when done right, is a magical collaborative experience: the user feels like the AI is a partner in creativity, offering ideas they may never have thought of, whether for fantasy adventures, sci-fi epics, or just silly bedtime stories.
- **Interactive Games and Role-Playing (AI as Game Master or NPC):** Beyond text adventures, LLMs are being used as **dungeon masters in tabletop**

**RPGs** or as dynamic NPCs (non-player characters) in video games. Imagine a Skyrim-like game where every villager or dragon can chat with you meaningfully, not just repeat pre-written lines. Companies and indie devs are exploring plugging LLMs into game engines so characters can respond to novel player questions. One striking example: using an LLM as an “AI *Dungeon Master*” for Dungeons & Dragons. This AI DM can **generate vivid descriptions, adapt the story to player choices in real-time, and ensure no two campaigns are alike**. It’s essentially an infinitely imaginative game master that never needs a break. Designing this involves feeding the game rules and world lore into the prompt so the AI doesn’t violate them (e.g., it should know that in this fantasy world, elves live longer than humans, or that magic has certain limits). The persona of the AI DM might be neutral narrator, or perhaps take on the flavor of the game (a spooky tone for a horror campaign). Feedback loops are super important here: players will naturally say “Actually, let’s ignore that quest and do X instead.” The AI must adapt, which is a design shift from traditional games where ignoring the main quest might lead to boring dead ends. With an LLM, ideally it will invent new content on the fly for the players’ creative detours. However, constraints and error handling are needed to avoid the “flock of geese” problem – i.e., the AI generating wild nonsense that breaks the game immersion (like the artifact that suddenly turned into geese in the Wayline blog example!). So designers will employ prompt constraints and possibly a validation step (making sure the AI’s output doesn’t break a list of don’ts, like “don’t suddenly kill the player’s character without reason,” etc.). This area is still experimental, but if cracked, it could revolutionize gaming – every playthrough truly unique and interactive on a conversational level.

- **Collaborative Ideation and Brainstorming Tools:** Another exciting use of conversational LLMs is as *ideation partners*. Think of a virtual brainstorming buddy available 24/7. For example, a team of designers might use an LLM-based chat tool to spitball concepts: “We need ideas for a sustainable packaging design”, and the AI could generate a list of creative approaches. LLMs are great at **expanding on a seed idea and offering many variations**, which can really help overcome creative block. When building such a tool, the conversation design might involve the AI asking questions to clarify the challenge (“Who is the target customer for these packages?”) and then using techniques like **divergent thinking prompts** (“Let’s list out 10 wild ideas, even if they seem impractical.”). The interface could allow users to refine or rate the ideas, which the AI then uses to hone further suggestions

(a feedback loop). One crucial design aspect is **avoiding bias and curation** – LLMs might sometimes over-index on common ideas, so designers could prompt it with “think of ideas *no one has tried before*” or combine it with random word prompts to spark novelty. Collaborative ideation bots can also be used in writing (co-writing a script or essay), coding (exploring different solution approaches), or business strategy (the AI acting as a sounding board for a startup pitch, asking the user the kind of tough questions an investor might). The benefit is an endlessly patient, idea-rich partner. However, UX-wise, one must ensure it doesn’t dominate the creative process – it should assist *your* creativity, not replace it. So the best designs keep the user in the driver’s seat: the user can say “Hmm, not quite, maybe more focus on eco-friendly materials,” and the AI adapts, much like a colleague might in a brainstorming meeting. When well integrated, an LLM can indeed function like an “AI collaborator,” boosting human creativity by contributing fresh perspectives on demand .

- **Educational and Coaching Bots:** We also see conversational LLMs being used as *tutors, career coaches, language practice partners*, and more. For instance, a student can chat with an AI that explains calculus concepts and answers questions, mimicking the dialogue one might have with a human tutor. The design here requires the AI to be pedagogically savvy – e.g., it should **comprehend the user’s intent and knowledge gaps** (did the student not understand a step in the solution? Are they asking a follow-up that indicates a misconception?) . The system prompt might enforce that the AI never just gives the answer outright to homework problems, but instead guides the student (to avoid cheating and encourage learning). Similarly, a language practice bot might speak in the target language, then switch to the user’s language to explain mistakes. The *interaction design* has to make the learning process engaging: perhaps using a bit of gamification (“Great job! 5 in a row correct!”) or adapting to the user’s skill level (simpler vocab for beginners). LLMs can also simulate interviewers for job practice, personal trainers for exercise motivation, or even just a knowledgeable friend to have debates with on various topics. Each of these scenarios needs careful tuning of persona and content to be effective and appropriate. But the common thread is that conversational AI can personalize and scale education/coaching in a way that static apps or documents can’t – it’s interactive and can respond to the learner’s specific questions in real time.

These examples just scratch the surface. What they highlight is that **LLM-driven conversational interfaces can shape user experiences in unexpected ways** – often blending utility with a bit of fun or human-like engagement. As designers in this space, we have to be creative and forward-thinking: we're not just solving old problems, we're envisioning entirely new modes of interaction. Whether it's giving someone comfort in the middle of the night via a chat, or collaboratively drafting a story, or managing a smart home through a friendly dialogue, the possibilities are vast.

However, all this possibility comes with responsibilities: ethical design (to avoid misuse or harm), inclusivity (making sure these interfaces work for people of different languages, abilities, etc.), and maintaining a **human-centered focus**. The best conversational interfaces empower and delight users, rather than confuse or replace them. As Bill Gates noted, AI interfaces might become as fundamental as the GUI did – meaning today's design students could be pioneers crafting the conventions everyone will use tomorrow.

In conclusion, designing conversations with LLMs is a multidisciplinary craft. It mixes the old (linguistics, HCI principles) with the new (prompt engineering, AI idiosyncrasies). It requires us to think about how **experiences** can be delivered through words and dialogue rather than pixels alone. It's an area filled with challenges (unpredictable AI behavior, anyone?) but also immense potential to create more natural, intuitive interactions with technology. By understanding the theories of interaction design and learning the specific techniques to tame and guide LLMs, we can create conversational interfaces that are not only useful, but meaningful and even magical for users. The conversation has only just begun – and you, as future designers and thinkers, will have a big role in shaping how those AI-human dialogues play out!

### Sources:

## References

- Cobus Greyling. (2023, June 19). *Prompt engineering, OpenAI & modes*. Medium. <https://medium.com/@cobusgreyling/prompt-engineering-openai-modes-8b7d83c52813>
- Greyling, C. (2023, July 25). *Imprinting personality on conversational AI*. Raga.ai. <https://www.raga.ai/blog/imprinting-personality-on-conversational-ai>

- Master of Code Global. (2023, March 16). *Conversation design in the age of LLMs*. Master of Code. <https://masterofcode.com/blog/conversation-design-in-the-age-of-llms>
- Microsoft. (n.d.). *System message*. Azure OpenAI Service documentation. <https://learn.microsoft.com/en-us/azure/ai-services/openai/how-to/system-message>
- Nielsen Norman Group. (2023, March 6). *Prompt structure in conversations with generative AI*. <https://www.nngroup.com/articles/prompt-structure-generative-ai/>
- Onlim. (2023, April 3). *The history of chatbots – From ELIZA to ChatGPT*. <https://onlim.com/en/history-of-chatbots-from-eliza-to-chatgpt/>
- RAGA.AI. (2023, July 18). *Memory and context awareness in LLM-based assistants*. <https://www.raga.ai/blog/memory-and-context-awareness>
- Syal, A. (2023, August 16). *Chatbots that think: Designing decision-making AI assistants*. Medium. <https://medium.com/@anirudhsyal/chatbots-that-think-7e19ac85d16e>
- Wayline Games. (2023, April 27). *Lessons from 10,000 hours of AI Dungeon Mastering*. <https://blog.wayline.games/lessons-from-10000-hours-of-ai-dungeon-mastering/>
- Wikipedia contributors. (2023, August 18). *AI Dungeon*. Wikipedia. [https://en.wikipedia.org/wiki/AI\\_Dungeon](https://en.wikipedia.org/wiki/AI_Dungeon)
- Wu, S. (2023, April 21). *Designing for AI beyond conversational interfaces*. Smashing Magazine. <https://www.smashingmagazine.com/2023/04/designing-ai-beyond-conversational-interfaces/>
- XLSCOUT. (2023, May 9). *Infinite possibilities: Enhanced brainstorming with LLMs*. <https://xlscout.ai/infinite-possibilities-enhanced-brainstorming-with-llms/>
- Zhou, M., & Gaffney, H. (2023). *Generative AI chatbots for mental health: Hope or hype?* npj Mental Health Research, 2(1), 14. <https://doi.org/10.1038/s44184-023-00014-7>